

Homework 3 - Pedestrian Detection using Histogram of Oriented Gradients

Topics:

- Data Preparation
- Feature extraction: Histogram of Oriented Gradients
- Classification: Support Vector Machine
- Object Detection: Sliding Window + Non Maximum Supression

Sources & Material

- Datasets:
 - https://www.cis.upenn.edu/~jshi/ped_html/
 - <http://pascal.inrialpes.fr/data/human/>
- Dataset preprocessed: <https://github.com/RashadGarayev/PersonDetection>
- Paper: Dalal and Triggs, "Histograms of oriented gradients for human detection," CVPR05
 - <https://hal.inria.fr/inria-00548512/document/>

Further Reading

- <https://www.pyimagesearch.com/2015/11/09/pedestrian-detection-opencv/>

1. Import Libraries

- OpenCV
- scikit-image
- scikit-learn

```
In [57]: import cv2 as cv
import matplotlib.pyplot as plt
from skimage.feature import hog
from skimage import exposure
import os
from sklearn import svm
import numpy as np
from sklearn.metrics import accuracy_score, confusion_matrix, average_precision_score, f1_score, recall_score,
import utils
```

2. Prepare Data

1. Create a dataset by loading all images
 - positive samples that contain pedestrians
 - negative samples that contain background objects
2. Split data into training and test data
 - split positive examples
 - split negative examples

```
In [58]: # Read Image files from local file system
positive_files = utils.ls("dataset/positive/")
negative_files = utils.ls("dataset/negative/")

split_pos = int(0.7*len(positive_files))
split_neg = int(0.7*len(negative_files))

train_positive_files, test_positive_files = positive_files[:split_pos], positive_files[split_pos:]
train_negative_files, test_negative_files = negative_files[:split_neg], negative_files[split_neg:]

# TODO:
# Read the images and store them in their corresponding lists

train_positive = [cv.imread(train_positive_files[i]) for i in range(len(train_positive_files)) ]
train_negative = [cv.imread(train_negative_files[i]) for i in range(len(train_negative_files)) ]

test_positive = [cv.imread(test_positive_files[i]) for i in range(len(test_positive_files)) ]
test_negative = [cv.imread(test_negative_files[i]) for i in range(len(test_negative_files)) ]
```

```
In [59]: # You can visualize the examples:
utils.showImages(train_positive[:6])
utils.showImages(train_negative[:6])
```



3. Feature Extraction

1. Apply Histogram of Oriented Gradients on all data
2. Create Training Labels

HOG - Example

```
In [60]: from skimage import data

img = data.astronaut()

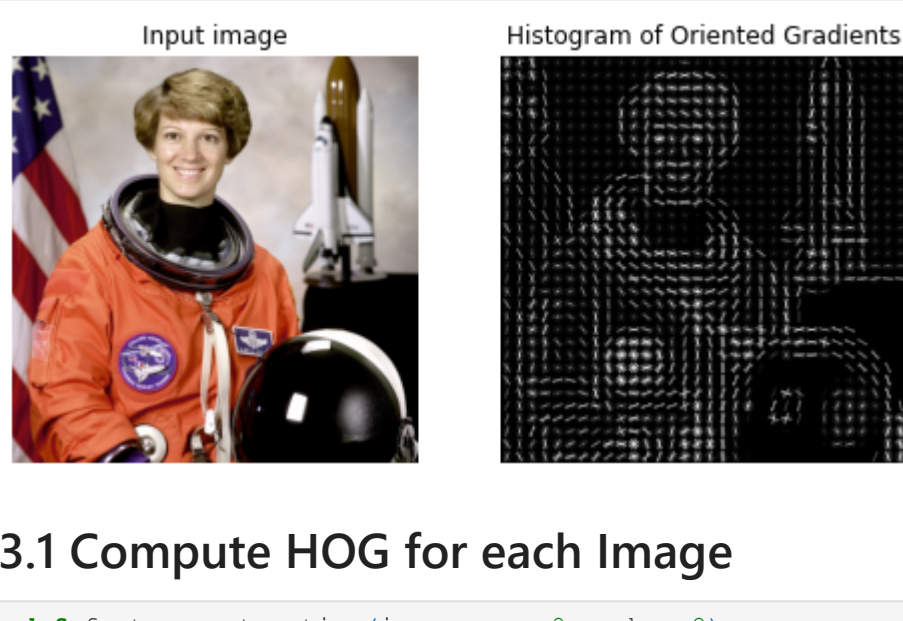
fd, hog_image = hog(img, orientations=8, pixels_per_cell=(16, 16),
                    cells_per_block=(3, 3), visualize=True, multichannel=True)

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(8, 4), sharex=True, sharey=True)

ax1.axis('off')
ax1.imshow(img, cmap=plt.cm.gray)
ax1.set_title('Input image')

# Rescale histogram for better display
hog_image_rescaled = exposure.rescale_intensity(hog_image, in_range=(0, 10))

ax2.axis('off')
ax2.imshow(hog_image_rescaled, cmap=plt.cm.gray)
ax2.set_title('Histogram of Oriented Gradients')
plt.show()
```



3.1 Compute HOG for each Image

```
In [61]: def feature_extraction(img, ppc = 8, cpb = 2):
    """
    img: A coloured image
    ppc: Pixels per cell
    cpb: Cells per block

    return: HOG feature vector
    The
    """
    # TODO: compute HOG feature vector
    # - try different parameters
    # - the output should be a one-dimensional array
    fd, hog_image = hog( img, orientations=8, pixels_per_cell=(ppc, ppc), cells_per_block=(cpb, cpb), visualize
    return fd

# Training Samples
train_positive_hog = np.array([feature_extraction(img) for img in train_positive])
train_negative_hog = np.array([feature_extraction(img) for img in train_negative])

# Testing Samples
test_positive_hog = np.array([feature_extraction(img) for img in test_positive])
test_negative_hog = np.array([feature_extraction(img) for img in test_negative])
```

3.2 Labels

```
In [62]: # Combine positive and negative samples into data matrices
X = np.concatenate((train_positive_hog, train_negative_hog))
X_test = np.concatenate((test_positive_hog, test_negative_hog))

# Create array of corresponding labels
Y = np.zeros(len(X))
Y[:len(train_positive_hog)] = 0
Y[len(train_positive_hog):] = 1

Y_test = np.zeros(len(X_test))
Y_test[:len(test_positive_hog)] = 0
Y_test[len(test_positive_hog):] = 1
```

4. Classifier

1. Initialize Model
2. Train Model
3. Evaluate Model

4.1 & 4.2

```
In [63]: svc = svm.LinearSVC(C=0.6, max_iter=10000)
svc.fit(X, Y)
```

Out[63]: LinearSVC(C=0.6, max_iter=10000)

4.3 Evaluate Model

```
In [64]: Y_pred = svc.predict(X_test)
```

Confusion Matrix

$$C = \begin{pmatrix} \text{True Negatives} & \text{False Positives} \\ \text{False Negatives} & \text{True Positives} \end{pmatrix}$$

```
In [65]: confusion_matrix(Y_test, Y_pred)
```

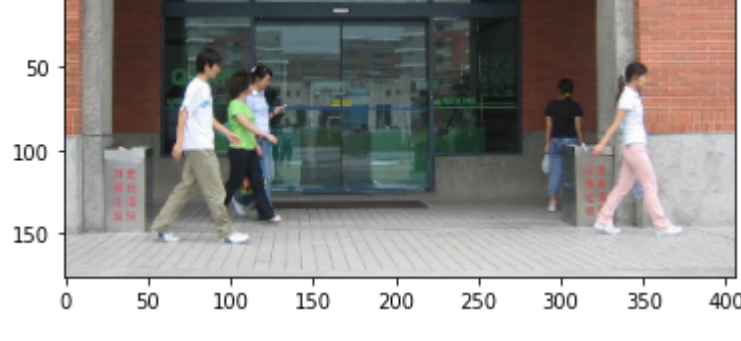
Out[65]: array([[651, 74],
 [60, 1176]], dtype=int64)

```
In [66]: print("AP:", average_precision_score(Y_test, Y_pred))
print("Accuracy:", accuracy_score(Y_test, Y_pred))
print("Recall:", recall_score(Y_test, Y_pred))
print("Precision:", precision_score(Y_test, Y_pred))
print("F1:", f1_score(Y_test, Y_pred))
```

AP: 0.925726731457598
Accuracy: 0.9316675165731769
Recall: 0.9514563106796117
Precision: 0.9408
F1: 0.9460981496379728

5. How to use the classifier on real images

```
In [67]: img = cv.imread("Images/FudanPed00036.png")[:, ::-1]
scale = 0.4
img = cv.resize(img, (int(scale*img.shape[1]), int(scale*img.shape[0])))
plt.imshow(img)
plt.show()
```



```
In [68]: height, width = 128, 64
import utils
def detect(img, svc, step = 10, height = 128, width = 64):
    """
    img: any image
    svc: Support Vector Classifier
    step: Step size of window operation

    return:
        result: a matrix with the predicted results [i,j] != 0 means that at the given window a pedestrian
        bounding_boxes [N, 4]: A matrix of bounding boxes corresponding to found pedestrians.
        bounding_boxes[i,0:2] describe the upper left corner and bounding_boxes[i,2:4] describe the bottom right
    """
    # Iterate over image and apply classifier
    result = np.zeros(img.shape[:2]) - 1
    for i in range(0, img.shape[0] - height, step):
        for j in range(0, img.shape[1] - width, step):
            result[i, j] = svc.predict(feature_extraction(img[i:i+height, j:j+width]).reshape(1, -1)) [0]

    y, x = np.nonzero(result == 0)
    bounding_boxes = np.stack([x, y, x+width, y+height], 1)
    return result, bounding_boxes

result, bounding_boxes = detect(img, svc)
```

```
In [69]: bounding_boxes_filtered = utils.non_max_suppression_fast(bounding_boxes, 0.5)
```

```
res = img.copy()
for i in range(len(bounding_boxes_filtered)):
    x1, y1 = bounding_boxes_filtered[i, 0], bounding_boxes_filtered[i, 1]
    x2, y2 = bounding_boxes_filtered[i, 2], bounding_boxes_filtered[i, 3]
    res = cv.rectangle(res.astype(np.float32), (x1, y1), (x2, y2), (0, 255, 0, 3).astype(np.uint8))

plt.imshow(res)
plt.show()
```




6 Deploy Model on Multiple Scales

The model only detects pedestrians in a specific resolution (64x128). Extend the application of the model to any resolution.

TODO

- Detect the images
- Scale people
- apply nonmaxima supression over multiple scales

```
In [70]: scales = [cv.GaussianBlur( img, (101,101), s ) for s in range(1,10)]
utils.showImages(scales)
```



```
In [71]: sub_ind = 1
plt.figure(figsize=(10,6))
for I in scales:
    result, bounding_boxes = detect(I, svc)
    bounding_boxes_filtered = utils.non_max_suppression_fast(bounding_boxes, 0.5)
    res = I.copy()
    for i in range(len(bounding_boxes_filtered)):
        x1, y1 = bounding_boxes_filtered[i, 0], bounding_boxes_filtered[i, 1]
        x2, y2 = bounding_boxes_filtered[i, 2], bounding_boxes_filtered[i, 3]
        res = cv.rectangle(res.astype(np.float32), (x1, y1), (x2, y2), (0, 255, 0, 3).astype(np.uint8))
    plt.subplot(3,3,sub_ind)
    plt.imshow(res)
    sub_ind += 1
```

